

DataCents Technical Paper

Title: A Machine Learning Framework for Predicting Loan Default Risk in Peer-to-Peer Lending

Authors: Noorelsalam Almakki, Madiha Maikzada, Myint Myat Zaw, Ahmed Hussein, Al-Hassan Sabeeh, Dadi Ishimwe

Date: August 8, 2025

Version: 1.0

Abstract

Peer-to-peer (P2P) lending platforms have democratized access to credit, but face the persistent challenge of managing borrower default risk. This paper presents the **DataCents** project, a data-driven initiative to identify the key predictors of loan default using a large-scale dataset from Lending Club. We employ a comparative approach, evaluating Logistic Regression, Random Forest, and XGBoost models to predict loan outcomes. Our analysis of over two million loan records reveals that the XGBoost model provides the most robust performance, achieving a ROC AUC of 0.72. Feature importance analysis, supported by SHAP values, identifies interest rate, loan grade, and FICO score as primary predictors. We use these findings to propose the **DataCents Risk Score**, a composite metric designed to provide a more nuanced and actionable assessment of credit risk for investors and platforms.

1. Introduction

The rise of the digital economy has led to the proliferation of peer-to-peer (P2P) lending platforms, which serve as intermediaries between individual lenders and borrowers. While these platforms offer greater efficiency and accessibility compared to traditional financial institutions, they are inherently exposed to credit risk—the risk of loss resulting from a borrower's failure to repay a loan. Inaccurate risk assessment can lead to significant losses for investors, erode platform trust, and limit the availability of capital.

Conventional credit scoring models, while valuable, may not capture the full spectrum of behaviors and attributes that signal default risk in the P2P ecosystem. This project addresses the following research question:

What are the key borrower and loan characteristics that best predict default risk in peer-to-peer (P2P) lending platforms in the United States?

To answer this, we leverage machine learning to analyze a comprehensive dataset of historical loan data, aiming to build a predictive model that is both accurate and interpretable.

2. Data and Methodology

2.1. Dataset

Our primary data source is the publicly available Lending Club loan dataset, spanning from 2007 to 2018. This dataset contains over 2.2 million loan records with more than 150 features per loan, including:

- **Borrower Information:** Employment length, annual income, FICO score range, home ownership status, and state.
- **Loan Characteristics:** Loan amount, term (36 or 60 months), interest rate, loan grade/sub-grade, and stated purpose (e.g., debt consolidation, credit card).
- **Credit History:** Date of earliest credit line, number of open credit accounts, and history of delinquencies.
- **Outcome Variable:** The loan_status field, which we mapped to a binary target: 1 for defaulted loans ("Charged Off") and 0 for fully paid loans.

2.2. Data Preprocessing

The raw data required significant cleaning and preparation. Our pipeline included the following steps:

1. **Filtering:** We retained only loans with a definitive status of "Fully Paid" or "Charged Off" to ensure a clear outcome for supervised learning.
2. **Handling Missing Values:** Features with a high percentage of missing values (>30%) were dropped. For others, we employed mean, median, or mode imputation where appropriate.
3. **Data Transformation:** Categorical features were one-hot encoded, and date features were converted into numerical representations (e.g., credit history length in months).

2.3. Feature Engineering

To enhance the predictive power of our models, we engineered several new features, including:

- **FICO Score:** We took the average of the fico_range_low and fico_range_high to create a single fico_score feature.
- **Credit History Length:** Calculated the number of years between the earliest_cr_line and the loan issue date.
- **State-Level Risk:** We incorporated a geographic risk factor by analyzing default rates on a per-state basis, creating a feature that represents the historical risk associated with a borrower's location.

2.4. Modeling Approach

We framed the problem as a binary classification task. The dataset was split into training (80%) and testing (20%) sets. Given the significant class imbalance (defaults represent a minority class), we used class weight adjustments in our models to penalize misclassifications of the minority class more heavily.

We evaluated three distinct models:

1. **Logistic Regression:** A linear model used as a baseline for performance.
2. **Random Forest:** An ensemble of decision trees to capture non-linear relationships.
3. **XGBoost (Extreme Gradient Boosting):** A powerful gradient boosting framework known for its performance and regularization capabilities.

3. Results and Analysis

3.1. Model Performance

The models were evaluated on the test set using ROC AUC, Precision, Recall, and F1-Score for the default class. XGBoost emerged as the superior model, offering the best balance between identifying true defaults (Recall) and maintaining reasonable precision.

Model	ROC AUC	Precision (Default)	Recall (Default)	F1-Score (Default)
XGBoost	0.72	0.32	0.67	0.44
Logistic Regression	0.71	0.31	0.67	0.43
Random Forest	0.71	0.55	0.06	0.10

3.2. Key Predictors of Default

Feature importance analysis from our XGBoost model highlighted several critical predictors of default risk. The most influential features, in descending order of importance, are:

- **Interest Rate (int_rate):** The most significant predictor. Higher rates are strongly correlated with higher default probability.
- **Loan Grade and Sub-Grade:** The platform's internal risk assessment is a powerful predictor.
- **FICO Score (fico_score):** Lower FICO scores are a primary indicator of credit risk.
- **Debt-to-Income Ratio (dti):** Borrowers with a higher proportion of income dedicated to existing debt are more likely to default.
- **Annual Income (annual_inc):** Lower income is associated with higher default risk.

3.3. Model Interpretability with SHAP

To ensure our model is not a "black box," we employed SHAP (SHapley Additive exPlanations). The SHAP summary plot provides a transparent view into the model's decision-making process, illustrating how each feature contributes to the prediction for individual loans. For instance, a high interest rate (red dot) pushes the SHAP value higher (towards default), while a high FICO score (low feature value, blue dot) pushes it lower (towards repayment).

4. The DataCents Risk Score

Based on our findings, we propose the conceptual framework for a **DataCents Risk Score**. This is a composite score that provides a holistic view of default risk by combining the most predictive features identified by our model. The score would be a weighted function of:

$$\text{RiskScore} = f(\text{int_rate}, \text{loan_grade}, \text{fico_score}, \text{dti}, \text{state_risk}, \dots)$$

This score would allow investors to quickly assess a loan's risk profile beyond traditional metrics.

5. Discussion and Limitations

5.1. Implications of Findings

Our results confirm that machine learning models can effectively identify complex patterns of risk in P2P lending data. The DataCents model provides a framework that can be used to augment existing credit assessment processes, leading to more profitable portfolios for investors and a more stable ecosystem for platforms.

5.2. Limitations

We acknowledge the following limitations:

- **Temporal Relevance:** The dataset ends in 2018 and does not capture macroeconomic shifts from the COVID-19 pandemic.
- **Class Imbalance:** Despite mitigation techniques, predicting a rare event like default remains challenging and can affect model precision.

5.3. Future Work

Future research will focus on:

- **Time-Series Analysis:** Modeling loan performance over time to predict default earlier in the loan lifecycle.
- **Bias and Fairness Audit:** Investigating the model for any demographic biases to ensure equitable outcomes.
- **Deployment:** Developing a user-facing dashboard or API for real-time risk assessment.

6. Conclusion

The DataCents project successfully demonstrates the utility of machine learning in decoding credit risk in P2P lending. Our XGBoost model provides a powerful and interpretable tool for identifying high-risk loans. By operationalizing these findings into a tool like the DataCents Risk Score, investors and platforms can make smarter, data-driven decisions that enhance profitability and build a more resilient financial ecosystem.